



Improving binary diffing speed and accuracy using community detection and locality-sensitive hashing: an empirical study

Chariton Karamitas¹ · Athanasios Kehagias²

Received: 7 December 2021 / Accepted: 13 September 2022
© The Author(s) 2022

Abstract

Binary diffing is a commonly used technique for detecting *syntactic* and *semantic* similarities and/or differences between two programs' binary executables (*not* source code). Here we present *REveal*, a binary diffing application. *REveal* is based on the detection of *Function Call Graph* (FCG) approximate isomorphism and improves both speed and accuracy, mainly by the use of two techniques. First, we propose the use of hierarchical *Community Detection* (CD) in executables' FCGs, for the purpose of detecting groups of densely connected functions, thus partitioning them in smaller groups. Moreover, we use *Locality-Sensitive Hashing* (LSH) for further grouping of similar functions in hash buckets. Both techniques are used in a *divide-and-conquer* fashion to simplify the diffing process of the programs being compared, practically reducing it to diffing of their FCG communities and LSH buckets.

Keywords Binary diffing · Community detection · Locality-sensitive hashing · Divide-and-conquer

1 Introduction

Binary diffing is a commonly used technique for detecting *syntactic* and *semantic* similarities or, equivalently, differences, between two programs' binary executables, when access to their source code is not an option. In most real world applications, binary diffing involves using various reverse engineering techniques to recover information from the subject executables, representing the recovered knowledge as a series of labeled graphs and computing a form of bijection between the latter. Examples of such graphs include the *Function Call Graph* (FCG), the *Control Flow Graph* (CFG), the *Program Dependence Graph* (PDG) and others. The said bijection is usually some form of graph isomorphism and this often means computing the *Maximum Common Subgraph* or *Multiple Common Subgraphs* of the two compared graphs.

Generally, the graph isomorphism problem belongs to the *NP* class and, consequently, related problems can only be solved approximately, by compromising precision in favor of acceptable running times.

A common technique, implemented in many binary diffing tools, including our prototype *REveal*, is approaching the problem of computing the Multiple Common Subgraphs of compared FCGs using *function matching*. In function matching, functions, that uniquely stand out based on their characteristics, are first matched, effectively creating an initial partial mapping of vertices between the compared graphs. Following that, functions "related to" already matched functions are compared and matched in an attempt to maximize the aforementioned mapping. This can be done (a) by using the recovered program structure (e.g. by matching successors and predecessors in the FCG [14,15,17]) and/or (b) by exploiting the locality characteristics inherent in the program linking process performed by most compilers (e.g. by matching functions nearby in the program's address space [28]). When no more new vertices can be matched, the process ends. In the literature, the former phase is usually referred to as the *pre-filtering phase*, while the latter as the *propagation phase*.

This paper presents two techniques for improving both the speed and accuracy of the overall binary diffing process as described in the previous paragraph. First, we propose the use

✉ Chariton Karamitas
huku@census-labs.com
Athanasios Kehagias
kehagiat@auth.gr

¹ CENSUS S.A. Department of Electrical and Computer Engineering, Aristotle University of Thessaloniki, Thessaloniki, Greece

² Department of Electrical and Computer Engineering, Aristotle University of Thessaloniki, Thessaloniki, Greece

of hierarchical *Community Detection* (CD) in executables' FCGs, for the purpose of detecting groups of densely connected functions, thus partitioning them in smaller groups. Moreover, we propose the use of *Locality-Sensitive Hashing* (LSH) for further grouping of similar, in terms of features, functions in hash buckets. Effectively, CD and LSH can be seen as complementary forms of partitioning; the former groups functions with respect to how they relate to their neighbors, while the latter probabilistically groups functions based on a hashing scheme of their features' values. In practice, this is equivalent to a 2-stage partitioning scheme where functions are divided into smaller groups by two mutually independent processes. Both techniques are used in a *divide-and-conquer* fashion to simplify the diffing process of the programs being compared, practically reducing it to diffing of their FCG communities and LSH buckets.

Overall, in this paper we make the following contributions:

- We present a fast and reliable hierarchical CD algorithm, based on *Node Similarity-based Algorithm* (NSA) [9] and *Louvain* [3,16], for partitioning FCGs in non-overlapping clusters of densely-connected functions. To our knowledge, use of CD algorithms in binary diffing has not been previously studied.
- We examine further splitting of communities in hash buckets of similar, in terms of characteristics, functions using LSH and, more specifically, a form of *minhashing*. [4]
- Based on the previous constructs, we develop a divide-and-conquer algorithm for binary diffing, which results, not only in speed improvements, but also in increased matching precision. We prove this by comparing our algorithms, implemented in REveal, against other binary diffing tools of the public domain, namely Diaphora and YaDiff, as well as an older version of REveal [28].

The rest of this paper is organized as follows. Section 2 gives an overview of previous work in the fields of binary diffing, CD as well as LSH. Preliminaries on graph theory and the notations used in this paper are given in Sect. 3. Section 4 presents the overall proposed CD and LSH based binary diffing algorithm, while Sect. 5 delves into its CD aspects and Sect. 6 elaborates on its LSH implementation. Last but not least, in Sect. 7 we evaluate our algorithm against those implemented in other tools and we conclude in Sect. 8 with a summary of current work and an overview of future research directions.

2 Previous work

When two graphs are similar, it is highly likely that their community structure will be similar as well. Based on this

simple observation, many graph theoretical problems can be approached as follows:

1. decompose the subject graph in smaller, but more compact, structures,
2. look for local solutions in the decomposed components,
3. form a global solution by combining local solutions.

This divide-and-conquer methodology has been successfully applied to many graph problems including (relevantly to our problem) the approximate graph isomorphism problem [1,24,32]. Even though community detection has been extensively used for tasks like knowledge extraction and graph classification, in various scientific domains, little has been done on computing the *Maximum Common Subgraph* of two graphs.

The idea of comparing two networks, by gradually compressing them in more compact representations, is presented at [42]. Using spatial clustering, cluster alignment using *Earth Moving Distance* (EMD) and *Graph Convolutional Networks* (GCNs), the authors essentially perform hierarchical community detection, based on label propagation, and matching of communities between the compared graphs. The *Hierarchical Graph Matching Network* (HGMN), as it is termed, can be used on any network, including CFGs and FCGs recovered from binary executables, as long as each vertex in the network can be abstractly represented by a feature vector (otherwise a feature vector of all-ones can be used). The effectiveness of HGMN is showcased in a series of experiments, where it is proven that it outperforms other graph-theoretic, as well as learning-based, approaches both in accuracy and efficiency.

Surprisingly, public literature on hierarchical graph isomorphism is very limited, with most research focused on graph classification. For example, seminal work on malware clustering is presented at [31]. First, the authors define cost functions for primitive graph edit operations and propose the use of GED for computing the similarity between two FCGs. Two major clustering algorithms are evaluated, namely *k-medoids* and *DBSCAN*. The authors conclude that DBSCAN is more promising for the task of partitioning malware samples or, equivalently, their FCGs into a set of an unknown number of clusters.

In [30], authors classify Android malware applications in clusters using DBSCAN. When a new malware sample is received, 4-tuples, holding structural information of methods, are first extracted. The 4-tuples of the new sample are matched with those of the representative mean of each cluster. A similarity score is computed as the ratio of matched methods over the total number of methods and the sample is added in the cluster with the highest matching score.

Community detection in FCGs is demonstrated in [37]. The authors investigate the potential of extracting cohesion

information on classes of object-oriented programs based on method communities recovered using a variety of algorithms including Louvain. Despite the fact that [37] is focused on computing software quality metrics, its conclusions might as well be applied for binary diffing purposes, especially when type information is available in the binary executables under comparison.

Yet another interesting technique of detecting *components* in binary code is presented in [29]. The authors construct a *Decomposition Graph* which is, in fact, the union of three other graphs, namely the *Sequence Graph* (SG), the *Data Reference Graph* (DRG) and the FCG. Newman's community detection algorithm with modularity optimization is used to recover the resulting graph's community structure in a hierarchical approach.

LSH was first introduced in [25]. This paper discusses the problem of nearest neighbor search and proposes LSH as a means of dealing with the *Curse of Dimensionality* (i.e., the fact that as the number of dimensions of the input space grows, nearest neighbor search algorithms become less and less efficient, to the point where they do no better than brute-force linear search). The techniques in [25] have a query time complexity of $\mathcal{O}(dn^{\frac{1}{\epsilon}})$ (where n is the number of points, and d the dimensionality). This was later improved by Gionis et al. [22] resulting in a query time of $\mathcal{O}(dn^{\frac{1}{\epsilon+1}})$.

More recently, LSH was used for binary diffing purposes in [35]. Basic blocks are first represented as a bag-of-words and converted to a large vector whose length is equal to the length of the alphabet (i.e., the number of instruction types of a given computer architecture). Each element in the aforementioned vector holds the number of occurrences of the corresponding instruction. Then, an arbitrary number of random hyperplanes of the same dimensionality are created. For each basic block, the angle it forms with each of the aforementioned hyperplanes is computed and the sign of the cosine of the angle is saved as a single bit (e.g. 0 means negative cosine, 1 positive). The concatenation of these bits, is effectively a hash of a basic block, which is used to group similar blocks and speed-up the nearest-neighbor search.

Binary diffing [14,15,17] is actively used, by researchers in various scientific domains, for tasks like malware classification [6,7,23], patch analysis [26], plagiarism detection [33], propagation of profiling information [40] and others. Various high quality tools for performing binary diffing tasks can be found in the public domain with [2,11,43] being some of them. An interesting summary of the capabilities of many modern binary diffing solutions is given in [12]. In this paper we extend REveal [27,28], our prototype tool which is being actively developed and tested.

3 Preliminaries

In the following sections we use uppercase letters to represent sets, and lowercase letters to represent primitive objects. Furthermore, we use the notation $\lceil x \rceil_P$ to indicate that x is rounded up to the nearest prime number.

3.1 Graph theory

We represent *digraphs* (i.e., directed graphs) with the notation $G = \langle V, E \rangle$, where V (or $V(G)$) is the digraph's vertex set and $E \subseteq V \times V$ (or $E(G)$) is the digraph's edge set. In what follows we will usually say, for brevity, *graph* when we actually mean *digraph*. Given a vertex $v \in V$, we define the set of *successors* of v as $\text{succ}(v) = \{s \mid (v, s) \in E\}$ and the set of *predecessors* of v as $\text{pred}(v) = \{p \mid (p, v) \in E\}$. The set of all *neighbors* of v is represented as $\text{neigh}(v) = \text{succ}(v) \cup \text{pred}(v)$.

We assume that, $\text{succ}(v)$, $\text{pred}(v)$ and $\text{neigh}(v)$ hold the corresponding vertex sets ordered by an arbitrary total ordering relation defined over V i.e., if $\text{succ}(v) = \{s_0, s_1, \dots, s_{n-1}\}$, for some $v \in V$, then $r(s_0) \leq r(s_1) \leq \dots \leq r(s_{n-1})$, where $r : V \rightarrow \mathbb{R}$ is a *sorting function* defined over the vertex set, that returns vertex *ordinals*; numbers via which the aforementioned total ordering is achieved. If, for a given sorting function r , $r(s_0) < r(s_1) < \dots < r(s_{n-1})$ holds, then r is said to be *strict*. Notice that from the previous definition, it follows that $r(s_i) \neq r(s_j) \forall i, j$ and $i \neq j$, that is, no two vertices should have the same ordinal number. Ordinals can be assigned to vertices in an application-specific manner.

We allow vertices and edges of a graph to carry *attributes* which can take arbitrary values. We use the notation $v.\text{attr}$ to refer to the value of attribute *attr* of vertex v and $(u, v).\text{attr}$ to refer to the attribute of edge (u, v) .

For each graph G , we define a *feature function* $f : V \rightarrow \mathbb{R}^d$ mapping each vertex of G to a feature vector of dimension d . Feature vectors are assumed to hold the characteristics of the objects represented by the corresponding vertices. When domain-specific features are not available, simple centrality metrics might be used (e.g. a feature vector of one element holding the vertex's degree). Similarly, we define the inverse of f , $f^{-1} : \mathbb{R}^d \rightarrow \mathcal{P}(V)$, which returns the set of all vertices of G that are characterized by a given feature vector.

Given a graph $G = \langle V, E \rangle$, we can partition vertices of G in non-overlapping communities $C_i \subseteq V$, $0 \leq i < N_C$, where N_C is the number of communities and $C_i \cap C_j = \emptyset$, $\forall i, j$ with $i \neq j$. The process via which communities are detected is referred to as *Community Detection* (CD). Furthermore, we can define the *Community Graph* (CG) to

be a graph whose set of vertices is $V = \{C_i \mid 0 \leq i < N_C\}$ and there is an edge from C_i to C_j iff $\{succ(v) \mid v \in C_i\} \cap \{pred(v) \mid v \in C_j\} \neq \emptyset$ (i.e., there's at least one edge in G from a vertex in C_i to a vertex in C_j and this holds for $i = j$ as well). Hierarchical CD algorithms repeat the same process on the resulting CG, thus producing another, higher level CG where vertices represent communities of communities. This process can be repeated N_L times, to generate a set of communities C_{li} , where $0 \leq l < N_L$ is the level index and i the index of a community in that level. In this case N_L CGs can be computed, one for each level. Notice that, usually, levels have different number of communities. More specifically, the higher the level, the less the number of communities, otherwise hierarchical community detection would make no sense.

3.2 Program representations

We define a program P to be a set of N_F functions $P = \{F_i \mid 0 \leq i < N_F\}$, a function F to be a set of N_B basic blocks (straight-line machine code sequences with no branches in, except to the entry, and no branches out, except at the exit) $F = \{B_i \mid 0 \leq i < N_B\}$ and a basic block B to be a set of N_I instructions $B = \{I_i \mid 0 \leq i < N_I\}$. Binary diffing involves comparing two programs, namely P_1 , the primary subject, and P_2 , the secondary subject, and forming a 1-1 mapping M , that corresponds functions of P_1 to functions of P_2 , $M = \{F_1 \rightarrow F_2 \mid F_1 \in P_1, F_2 \in P_2\}$. Similar mappings can be further created for function basic blocks and instructions, depending on the required level of detail.

For each function F , in a program P , we define a digraph $CFG = \langle V, E \rangle$, where V is the set of F 's basic blocks, while the set of edges E denotes the possible execution flow paths between the function's basic blocks. If $e = (B_i, B_j) \in E$, then control flow can reach basic block B_j immediately after B_i (i.e., $B_j \in succ(B_i)$). Digraph CFG is usually referred to as F 's *Control Flow Graph* (CFG).¹

A program P can also be treated as a digraph of program functions, referred to as the program's *Function Call Graph* (FCG). FCG of program P is a digraph whose vertices correspond to individual functions (or, equivalently of function CFGs), that is $V = \{F \mid F \in P\}$. For each control transfer instruction in basic block $B_k \in F_i$, that transfers execution to basic block $B_l \in F_j$ (this also covers the case of $F_i = F_j$), an edge (F_i, F_j) exists in E . Multiple such edges (F_i, F_j) can be replaced with a single weighted edge with a weight equal to the number of calls from F_i to F_j .

As uniquely identifying vertices by centrality measures is generally not possible, a common ordering of vertices in program related graphs (like FCGs and CFGs) can be

based on the address, in program memory, of the structure represented by each vertex (like functions or basic block accordingly). That is r , defined in the previous section, is $r_{address} : V \rightarrow \mathbb{N}$, where $r_{address}(v)$ is a function that returns the address of vertex v in program memory. For CGs, an address can be assigned to each vertex based on the minimum or the maximum address of the elements contained in each community.

Furthermore, we define a *dataset* as $D = \langle G, S, r, f \rangle$, where G is a CG (i.e., community graph) or FCG, $S \subseteq V(G)$, r a sorting function defined over $V(S)$ and f the feature function of G . Generally, the presented algorithms process functions of P in datasets using various strategies. For example, given a function $F \in P$, a dataset can be constructed consisting of P 's FCG, $S = neigh(F)$, $r_{address}$ and the FCG's feature function. When G is a FCG, the corresponding dataset is referred to as a *function dataset*, in the opposite case it is termed a *community dataset*.

Last but not least, for a previous discussion on feature selection for vertices of a FCG the reader is referred to [27,28] and [39].

3.3 Locality sensitive hashing

Locality Sensitive Hashing (LSH) is a process by which similar inputs are hashed to the same bucket. Unlike normal cryptographic hash functions, that need to minimize the number of hash collisions, LSH functions attempt to maximize them for inputs which are "close to each other" in a given metric space \mathcal{M} . More specifically, given two points $p, q \in \mathcal{M}$ and a LSH function $h : \mathcal{M} \rightarrow \mathbb{N}$, mapping points in \mathcal{M} to hash bucket indexes, the following should hold:

1. If $d(p, q) < R$, then $h(p) = h(q)$ with probability at least p_1
2. If $d(p, q) \geq cR$, then $h(p) = h(q)$ with probability no more than $p_2 \ll p_1$

In the above relations, R is a distance threshold, below which points are considered to be "close by", while $c > 1$ is the approximation factor. The first relation indicates that, when p and q are nearby, according to a distance function d defined on \mathcal{M} 's points, then they should hash to the same bucket with probability at least p_1 . The second relation represents the case of a hash collision for inputs which are not nearby. Obviously, for the LSH function h to be meaningful, p_1 must be much higher than p_2 , so that hashing dissimilar objects in the same bucket is less probable than the opposite case.

Among its many applications, LSH is commonly used for solving the k -nearest neighbor problem. A family of L LSH functions $H = \{h_i \mid 0 \leq i < L\}$ (e.g. *minhashing* [4], based on min-wise independent permutations, or *simhashing* [8],

¹ Even though the CFG (and the FCG defined in the sequel) are *digraphs*, we will follow standard usage and call them *graphs*.

based on random projections) is first constructed. Inputs are hashed using all such functions and eventually are split up in buckets. Given an input query point q , its k -nearest neighbors can be found by hashing q using all hash functions in H and looking for neighbors in the resulting L buckets. Search stops when k candidates have been found, or k “best” points have been drawn from the overall result set.

4 A community-based function matching algorithm

We begin by giving a high level overview of the overall matching process, carried out by REveal, in Algorithm 1 and a call-graph of the most essential procedures, presented in this paper, in Fig. 1.

Before proceeding to a detailed description of the above blocks (in separate subsections), let us give a brief overview of the workflow described in Fig. 1.

- At the top of the diagram we have `match_programs` which, upon conclusion, will return the mapping MF of functions of program P_1 to functions of program P_2 . This function calls `match_comms` and `match_funcs`.
- `match_comms` is called first to detect the community structure of programs P_1 and P_2 , effectively aggregating functions into groups (i.e. communities). It returns a mapping MC of communities of program P_1 to communities of program P_2 .
- Given MC , populated by `match_comms`, `match_funcs` begins the actual function matching process. It basically executes consecutive rounds of exact and inexact matching of functions within the matched communities (in MC) until no more new matches can be found.
- `match_funcs_in_comms` is the actual workhorse used by `match_funcs`. It is able to perform rounds of exact and inexact function matching by utilizing the information in MC . A boolean argument passed to `match_funcs_in_comms` determines whether exact or inexact matching will take place.
- `match`, called by `match_comms` and `match_funcs_in_comms`, matches datasets (as defined in Sect. 3.2), in an abstract manner. It expects two datasets as inputs and uses the data features to match dataset elements. It does not distinguish between community datasets and function datasets, as the underlying matching strategies and algorithms are the same. It keeps applying various matching strategies (described below) until new matches cannot be found.
- Blocks named `*_matcher` implement the actual matching strategies as described in [28]. Matching strategies are, practically, ways of constructing “smaller” datasets

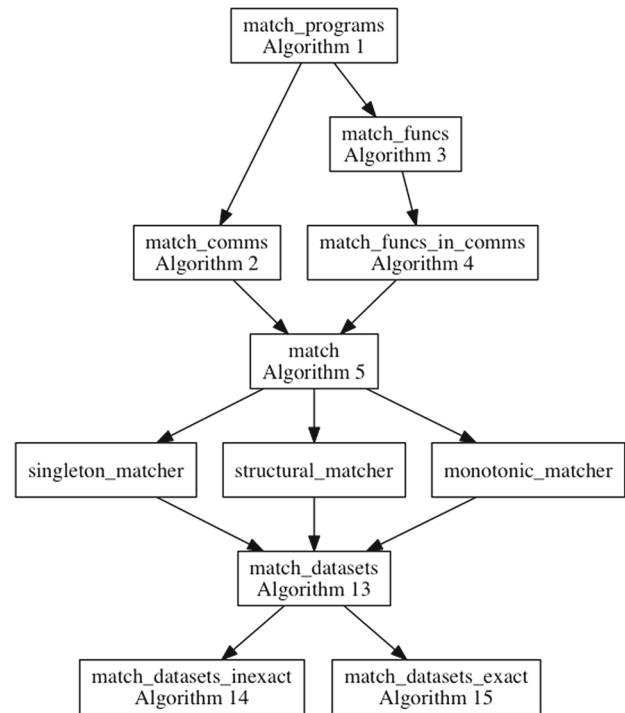


Fig. 1 Call-graph of most essential algorithms presented in this paper

of related functions, so that those smaller datasets can be matched for speed and efficiency purposes. For example, `structural_matcher` will construct temporary datasets, consisting of neighbors of already matched functions, and will attempt to match these to expand MF with new entries.

- `match_datasets` takes these smaller datasets, generated by the abovementioned blocks, and calls `match_datasets_inexact` or `match_datasets_exact` to actually populate a mapping with new entries. The said mapping can either be MC or MF since, as it was previously mentioned, the same abstract algorithms are used for both community and function matching.

Before proceeding, it is also useful to detail the ways in which the current implementation of REveal improves upon the previous one [27,28]. Briefly, the new version of REveal:

- Implements community detection and matching using the algorithms presented in Sect. 5.
- Performs inexact matching, in addition to exact implemented in the previous version.
- Increases matching speed using LSH (Sect. 6).

We continue by elaborating on the various functional blocks of Fig. 1.

4.1 Algorithm 1: `match_programs`

Algorithm 1 expects the two programs under comparison, namely P_1 and P_2 , to be given as inputs. The output of Algorithm 1 is the set MF (stands for *Matched Functions*), with elements of the form $F_1 \rightarrow F_2$, mapping function F_1 of P_1 to function F_2 of P_2 .

Algorithm 1 Main matching algorithm

```

1: procedure match_programs( $P_1, P_2$ )
2:    $D_1 \leftarrow \langle FCG_{P_1}, V(FCG_{P_1}), r_{address}, f_{P1} \rangle$ 
3:    $D_2 \leftarrow \langle FCG_{P_2}, V(FCG_{P_2}), r_{address}, f_{P2} \rangle$ 
4:    $MC \leftarrow \text{match\_comms}(D_1, D_2)$ 
5:    $MF \leftarrow \text{match\_funcs}(D_1, D_2, MC)$ 
6:   return  $MF$ 
7: end procedure

```

Given input programs P_1 and P_2 , `match_programs` begins by constructing the corresponding datasets D_1 and D_2 at lines 2 and 3. These datasets hold the two programs' FCGs, while the sorting functions are set to $r_{address}$ and the feature functions, f_{P1} and f_{P2} , return program function features (e.g. [27,28,39]) of P_1 and P_2 respectively. The algorithm continues by detecting and matching communities in the programs' FCGs, by means of `match_comms` (line 4), and proceeds by using the recovered community structure, in MC (stands for *Matched Communities*), in order to match program functions via `match_funcs` (line 5). Algorithm 1 terminates by returning MF .

4.2 Algorithm 2: `match_comms`

Algorithm 2, begins by initializing MC to an empty list (line 2). More specifically, MC is a list of sets, with the set at index i holding elements of the form $C_1 \rightarrow C_2$, where $C_1 \subseteq P_1$ and $C_2 \subseteq P_2$, representing matched communities at level $num_levels - 1 - i$ of the community hierarchy of the compared programs. The community hierarchies themselves are recovered at lines 3–4, using techniques and algorithms presented in Sect. 5, and implemented by `detect_comms`. The aforementioned function returns a list of community datasets, with the i -th set holding information on the community structure at level i in the corresponding program's community hierarchy. The higher the value of i the bigger (i.e. more abstract) the communities. For example, assuming P_1 consists of a 2-level community hierarchy, $CD_{1,0}$ represents communities of functions, while $CD_{1,1}$ communities consisting of $CD_{1,0}$'s communities. Evidently, the process of community detection is repeated for both programs and the results are stored in CD_1 and CD_2 for the first and second programs respectively.

Algorithm 2 Community detection and matching

```

1: procedure match_comms( $D_1, D_2$ )
2:    $MC \leftarrow \emptyset$ 
3:    $CD_1 \leftarrow \text{detect\_comms}(D_1)$ 
4:    $CD_2 \leftarrow \text{detect\_comms}(D_2)$ 
5:    $num\_levels \leftarrow \min(\|CD_1\|, \|CD_2\|)$ 
6:   for all  $i \leftarrow 1$  to  $num\_levels$  do
7:      $MC_i \leftarrow \emptyset$ 
8:      $change \leftarrow true$ 
9:     while  $change = true$  do
10:       $change \leftarrow false$ 
11:      if  $\text{match}(CD_{1,-i}, CD_{2,-i}, MC_i, false) > 0$  then
12:         $change \leftarrow true$ 
13:      end if
14:      if  $\text{match}(CD_{1,-i}, CD_{2,-i}, MC_i, true) > 0$  then
15:         $change \leftarrow true$ 
16:      end if
17:    end while
18:     $MC \leftarrow MC \cup \{MC_i\}$ 
19:  end for
20:  return  $MC$ 
21: end procedure

```

With a N_{L1} -level community hierarchy having been recovered from P_1 and a N_{L2} -level from P_2 , a process of matching the communities in each level begins, with only $\min(N_{L1}, N_{L2})$ (line 5) levels actually being considered. The process of matching communities is carried out level by level starting from higher ones (i.e. larger, more abstract communities) to lower ones (i.e. smaller, more concrete communities). Notice how the following loop iterates from 1 to num_levels (line 6), but indices in $CD_{1,-i}$ and $CD_{2,-i}$ are negative signifying reverse element access (i.e. -1 is the last set, -2 the second to last etc.). In each iteration a new empty set MC_i is initialized (line 7) and `match` (defined later on) is invoked to match communities in the community datasets. A round of exact matching (last argument of `match` is *false*) follows a round of inexact matching (last argument of `match` set to *true*) and the whole process repeats until no more communities can be matched (lines 8–17). In each iteration MC_i is appended in MC (line 18). Ultimately, MC is returned (line 20) for use by `match_funcs`.

4.3 Algorithm 3: `match_funcs`

The entry point to function matching can be seen in Algorithm 3. The two function datasets D_1 and D_2 , as well as the matched community hierarchy information in MC , returned by `match_comms`, are expected as inputs. MF , is a set that holds matched function pairs and is initially set to the empty set (line 2), while the loop, at lines 4–12, repeatedly appends new entries in it. At the core of the aforementioned loop, `match_funcs_in_comms` (defined later on) is called twice at lines 6 and 9. The last argument determines the type of function matching round that will take place; *false* means exact matching, while *true* inexact. In

Algorithm 3 Function matching

```

1: procedure match_funcs( $D_1, D_2, MC$ )
2:    $MF \leftarrow \emptyset$ 
3:    $change \leftarrow true$ 
4:   while  $change = true$  do
5:      $change \leftarrow false$ 
6:     if  $match\_funcs\_in\_comms(D_1, D_2,$ 
        $MC, MF, false) > 0$  then
7:        $change \leftarrow true$ 
8:     end if
9:     if  $match\_funcs\_in\_comms(D_1, D_2,$ 
        $MC, MF, true) > 0$  then
10:       $change \leftarrow true$ 
11:    end if
12:  end while
13:  return  $MF$ 
14: end procedure

```

essence, once no more functions can be matched exactly, inexact matches are looked up and if that generates new findings, the exact matching round needs to be repeated. When the loop breaks, MF is returned (line 13).

4.4 Algorithm 4: *match_funcs_in_comms*

Algorithm 4, which is at the heart of our divide-and-conquer approach, expects 5 input parameters; two function datasets D_1 and D_2 , the set of sets of matched communities, as returned by *match_comms* (defined previously), an initially empty set of matched functions MF and a boolean argument *is_inexact*, whose value determines whether the current matching round will be an exact or inexact one. The main logic of Algorithm 4 is implemented in two nested loops at lines 6–14 and 7–13. The outer loop iterates over all sets of matched communities (one for each level in the community

hierarchy). The inner loop iterates over all matched communities of each such set. Given a valid match $C_1 \rightarrow C_2$, two new function datasets D'_1 and D'_2 are created. The vertex subsets, in these datasets, are set to the functions contained in C_1 and C_2 respectively, effectively forming a smaller set of potential match candidates. It's essential to note that, since the outer loop, at lines 6–14, iterates through all community levels, all but the last level hold communities of communities. In this case, $\{F \mid F \in C\}$ practically means “all functions which belong to communities of community C transitively”. The new, smaller datasets are passed to function *match*, which populates MF with matched function pairs of the form $F_1 \rightarrow F_2$, where $F_1 \in C_1$ (and, of course, $F_1 \in P_1$) and $F_2 \in C_2$ (and $F_2 \in P_2$). Once all community levels have been handled, a final attempt to match functions in D_1 and D_2 is made at line 15. Note that D_1 and D_2 are the initial function datasets and that this step accounts for any potential discrepancies in the way communities were recovered, by giving a chance to functions, erroneously identified as belonging to different communities, to be matched if they are the same (exact matching case) or similar (inexact matching case). A third, outmost loop, at lines 4–18, guarantees that the overall process repeats as long as new findings are detected. Finally the total number of additional matched function pairs, appended in MF , is returned at line 19.

4.5 Algorithm 5: *match*

The pseudocode for method *match* is given in Algorithm 5. It is comprised of an outer loop (lines 4–15) that executes as long as *singleton_matcher*, *structural_matcher* or *monotonic_matcher* return an integer greater than 0. For brevity, the pseudocode and semantics of the three

Algorithm 4 Function matching

```

1: procedure match_funcs_in_comms( $D_1, D_2,$ 
    $MC, MF, is\_inexact$ )
2:    $n \leftarrow \|MF\|$ 
3:    $change \leftarrow true$ 
4:   while  $change = true$  do
5:      $change \leftarrow false$ 
6:     for all  $MC_i \in MC$  do
7:       for all  $C_1 \rightarrow C_2 \in MC_i$  do
8:          $D'_1 \leftarrow \langle G(D_1), \{F \mid F \in C_1\}, r_{address}, fp_1 \rangle$ 
9:          $D'_2 \leftarrow \langle G(D_2), \{F \mid F \in C_2\}, r_{address}, fp_2 \rangle$ 
10:        if  $match(D'_1, D'_2, MF, is\_inexact) > 0$  then
11:           $change \leftarrow true$ 
12:        end if
13:      end for
14:    end for
15:    if  $match(D_1, D_2, MF, is\_inexact) > 0$  then
16:       $change \leftarrow true$ 
17:    end if
18:  end while
19:  return  $\|MF\| - n$ 
20: end procedure

```

Algorithm 5 Match round

```

1: procedure match( $D_1, D_2, M, is\_inexact$ )
2:    $n \leftarrow \|M\|$ 
3:    $change \leftarrow true$ 
4:   while  $change = true$  do
5:      $change \leftarrow false$ 
6:     if  $singleton\_matcher(D_1, D_2, M,$ 
        $is\_inexact) > 0$  then
7:        $change \leftarrow true$ 
8:     end if
9:     if  $structural\_matcher(D_1, D_2, M,$ 
        $is\_inexact) > 0$  then
10:       $change \leftarrow true$ 
11:    end if
12:    if  $monotonic\_matcher(D_1, D_2, M,$ 
        $is\_inexact) > 0$  then
13:       $change \leftarrow true$ 
14:    end if
15:  end while
16:  return  $\|M\| - n$ 
17: end procedure

```

aforementioned methods are omitted (for more information readers are referred to [28]).

5 Community detection

5.1 Introductory remarks

Before we present in more detail the components of Algorithm 2 (`match_comms`), let us briefly discuss the community detection problem.

An important aspect of graphs is their *community structure* which can be recovered by *vertex clustering*. By this we mean the separation of vertices into clusters, with many edges joining vertices of the same cluster and comparatively few edges joining vertices of different clusters. The community detection problem is not well posed, because there are many possible definitions of cluster (i.e., community) and choosing a particular definition is to some degree arbitrary. Hence many different approaches appear in the related literature [19,20]. However it is worth noting that intuitively, we can recognize that many real-life graphs do exhibit community structure. That is, while detecting groups of densely-connected or related vertices may indeed be possible by visual inspection, no universal algorithm is currently known to reliably do this.

Choosing a community detection algorithm for graph analysis applications is a task that requires many parameters to be taken into account. Apart from an algorithm's ability to detect high quality communities, its capability to handle large networks, possibly with millions of vertices, and the memory footprint, execution time is also an essential aspect and a major criterion for making the final choice. One of the most popular community detection algorithms, is the *Louvain* algorithm [3,16], characterized by its capability to handle very large graphs with reasonable time and memory requirements. Hence we have decided to use this algorithm as a core component of our REveal algorithm, the reference implementation presented in this paper.

Despite its good performance, it is well known that the Louvain community detection algorithm suffers from some weaknesses. To begin with, according to [38], a vertex acting as a “bridge” between two groups of vertices in a community, may, for modularity optimization purposes, be moved in a different community, resulting in the former being disconnected. Iteratively applying Louvain may further worsen the problem. Furthermore, being a *Local Modularity Optimization* (LMO) algorithm, the Louvain algorithm may fail to discover small communities due to the *resolution limit* problem, described by Fortunato and Barthelemy [18]. To highlight another weakness of the Louvain algorithm consider the following: given two versions of the same program, community detection in their FCGs can be thought of as the

problem of detecting communities in a dynamic graph¹. In this case, due to its inherent randomness, the Louvain algorithm produces unstable partitionings, for t and $t + 1$, and is thus unsuitable for this purpose. This “instability” problem is further explored in [10].²

Due to the aforementioned problems, in certain applications ([21,34]), the Louvain algorithm is only used to improve the modularity of an existing partitioning, instead of generating one from scratch. That is, in the first iteration of the Louvain algorithm, graph vertices are already grouped in a set of preliminary initial communities, possibly produced by other means, as opposed to forming singleton communities. Then, Louvain is applied to fine-tune the initial partitioning. In [34], for example, the authors present a divisive community detection algorithm, referred to as MILPA, which uses the Louvain algorithm in such a post-processing fashion.

Following this approach, we propose the use of a *Label Propagation Algorithm* (LPA), for the purpose of computing an initial set of communities on a given FCG, followed by an application of the Louvain algorithm for improving the aforementioned community structure's modularity score. The LPA algorithm used is a modified and trimmed down version of *Node Similarity-based Algorithm* (NSA) described by Cheng et al [9]. The process is repeated as long as the number of detected communities changes. Our final community detection algorithm is summarized in Algorithm 6. As it will become apparent in the following sections, to “stabilize” the partitionings produced, the community detection algorithms were modified to introduce determinism in the overall process.

Let us now discuss, in separate subsections, the algorithms invoked by Algorithm 2 (`match_comms`).

5.2 Algorithm 6: `detect_comms`

Algorithm 6 expects a function dataset D as input, recovers the community hierarchy of D 's FCG and, finally, computes and returns a list of community datasets (one element per level of the community hierarchy). `detect_comms` begins by initializing CD (i.e. community datasets) to the empty list (line 2). CD is a list of sets, with the i -th element holding the community dataset at level i of the program's community hierarchy. Procedures `compute_vertex_ranks` (line 4) and `compute_edge_ranks` (line 5) are, then, called to assign vertex and edge *ranks* respectively to the FCG. The input dataset D , which holds the program FCG, is then converted to a new dataset (line 6) where FCG vertices are ordered by sorting function r_{rank} , which in turn sorts graph

² The dynamic graph in question, in time t , corresponds to the FCG of the first version, while, in time $t + 1$, to the FCG of the second version. This idea is based on the fact that the first FCG can be transformed into the second via a series of vertex and/or edge additions and/or deletions.

Algorithm 6 Community detection

```

1: procedure detect_comms( $D$ )
2:    $CD \leftarrow \emptyset$ 
3:    $G_0 \leftarrow G(D)$ 
4:   compute_vertex_ranks( $G_0$ )
5:   compute_edge_ranks( $G_0$ )
6:    $CD_0 \leftarrow \langle G_0, V(G_0), r_{rank}, f(D) \rangle$ 
7:    $C_0 \leftarrow \emptyset$ 
8:   for all  $i \leftarrow 1$  to ... do
9:      $C_i \leftarrow \text{detect\_initial\_comms}(CD_{i-1})$ 
10:     $C_i \leftarrow \text{improve\_comms}(CD_{i-1}, C_i)$ 
11:     $G_{i-1} \leftarrow G(CD_{i-1})$ 
12:     $G_i \leftarrow \text{compute\_induced\_graph}(G_{i-1}, C_i)$ 
13:     $CD_i \leftarrow \langle G_i, V(G_i), r_{rank}, f_{com} \rangle$ 
14:    if  $\|C_i\| == \|C_{i-1}\|$  then
15:      break
16:    end if
17:     $CD \leftarrow CD \cup \{CD_i\}$ 
18:  end for
19:  return  $CD$ 
20: end procedure

```

vertices according to the ranks assigned previously. The concept and rationale behind vertex ranking is further analyzed in Sect. 5.3. In this context, vertex ranking effectively results in functions, that process datasets that use r_{rank} , accessing vertices of the dataset's graph in order of increasing rank (e.g. in $\text{succ}(v)$, $\text{pred}(v)$ or $\text{neigh}(v)$ for a $v \in V(G)$).

Community hierarchy recovery is an iterative process (lines 8–17) which is repeated as long as the number of detected communities changes between levels. In each iteration, the modified NSA algorithm (*detect_initial_comms*, line 9) is applied on the current dataset CD_{i-1} (with D_0 being a function dataset holding the FCG, sorted using r_{rank}). The result is a set of communities, which is given as input to the directed Louvain algorithm [16] (*improve_comms*, line 10). The latter's task is to improve the quality of the partitioning, using local modularity optimization heuristics. Next, the induced graph of communities is computed (line 12). In this graph, vertices correspond to G_{i-1} 's communities and edges connect communities of neighboring vertices. More specifically, given G_{i-1} and its decomposition into communities C_i , *compute_induced_graph* constructs a new graph G_i with the following characteristics;

1. $V(G_i)$ is the set of communities of G_{i-1} ; $V(G_i) = \{C_0, C_1, \dots, C_{N_C}\}$
2. $C_i.rank = \sum_{v \in C_i} v.rank$ where $v \in V(G_{i-1})$, that is, vertex ranks of G_i are computed by summing the ranks of vertices in the corresponding community.
3. Edge ranks in the new graph are computed as the sum of the ranks of the original graph's edges as follows: $(C_i, C_j).rank = \sum_{u \in C_i, v \in C_j, (u,v) \in E(G_{i-1})} (u, v).rank$.

Last but not least, in each iteration, a new community dataset is created, to hold the results of the current partitioning (line 13) and is appended in the list of community datasets (line 17).

5.3 Vertex and edge ranking

The results produced by many agglomerative community detection algorithms, including [9], as well as [3], depend on the order that graph vertices are traversed. More specifically, in [3] the authors claim that preliminary results on several test cases seem to indicate that the ordering of the vertices does not have a significant influence on the modularity that is obtained. However, traversal order may impact performance, and more importantly, the structure of the communities produced as output. In [9] the authors traverse graph vertices in order of decreasing degree. Ties are broken by choosing an arbitrary vertex from a set of vertices with the same degree. Consequently, applying the algorithm twice, even on the same graph, might produce two different results.

For graph comparison purposes, the communities of the compared graphs, detected by whatever algorithm is used for that purpose, should be as similar as possible. Unlike other scientific domains, where detected communities are more valuable when community vertices are *semantically* related, for graph isomorphism related problems, communities do not have to make sense, as long as the community structures, recovered from the binaries under comparison, are as identical as possible. In an attempt to minimize the effect of randomness in community detection algorithms, a high quality ordering of graph vertices should be adopted.

Ideally the corresponding sorting function should be strict, but this is not easy in practice. Given two programs P_1 and P_2 , one needs to find a sorting function r , defined over $F_1 \cup F_2$, such that $r(F_1) = r(F_2)$, with $F_1 \in P_1$ and $F_2 \in P_2$, whenever F_1 and F_2 are the same function. This however has many problems in practice. First and foremost, it might be the case that F_1 and F_2 are byte-by-byte equal but are not the same function (e.g. this is quite common in binary code that makes extensive use of C++ templates). Furthermore, finding such an r requires that all mappings $F_1 \rightarrow F_2$ are known in advance, which is impossible, as this is the problem we are asked to solve in the first place. Consequently, only approximations can be made, that is we need a sorting function r that produces $r(F_1) = r(F_2)$ with a high probability if F_1 and F_2 are equal.

5.3.1 Algorithms 7–9: *compute_vertex_ranks* and associated algorithms

Our sorting function r is based on the notion of vertex ranks. Vertex ranks are integer values, and more specifically

prime numbers, assigned to vertices of the program FCG by `compute_vertex_ranks`, shown in Algorithm 7.

Algorithm 7 Compute initial vertex ranks

```

1: procedure compute_vertex_ranks(D)
2:    $G \leftarrow G(D)$ 
3:   for all  $v \in V(G)$  do
4:      $v.rank \leftarrow \lceil r_{inf}(v) \rceil_p$ 
5:   end for
6:   for all  $i \leftarrow 1$  to 16 do
7:     for all  $v \in V(G)$  do
8:        $w \leftarrow v.rank$ 
9:       for all  $u \in \text{neigh}(v)$  do
10:         $w \leftarrow w * u.rank$ 
11:      end for
12:       $v.rank' \leftarrow w$ 
13:    end for
14:    for all  $v \in V(G)$  do
15:       $v.rank = \lceil v.rank' \rceil_p$ 
16:      del  $v.rank'$ 
17:    end for
18:  end for
19: end procedure

```

Given a function dataset, Algorithm 7 computes initial rank values for each function in the dataset's FCG (lines 3–5) and then executes 16 Weisfeiler–Lehman [41] rounds (lines 6–18) to obtain a high quality coloring based on these ranks (16 was chosen empirically, so that each vertex' color is affected by that of neighbors at most 16 “steps” away). Ranks are, in fact, stored as vertex attributes and their initial value is set to the quantity returned by r_{inf} (defined later on) rounded up to the next prime number. Each Weisfeiler–Lehman round (lines 7–13) computes the new rank of each vertex based on the ranks of its neighbors, using plain integer product, which is guaranteed to be unique for vertices that have similarly ranked neighbors. The new rank value is stored in attribute $rank'$ and eventually replaces attribute $rank$ (lines 14–17) after being rounded up to the next prime number. Once all 16 iterations have completed, Algorithm 7 returns and $rank$ attributes hold a high quality coloring of the FCG vertices. An interesting observation is that the *for* loops at lines 3–5, 7–13 and 14–17 traverse the set of graph vertices in order of increasing address. Indeed, dataset D passed to `compute_vertex_ranks` in Algorithm 6 uses $r_{address}$ as a sorting function. Even though this fact alone is not pivotal for the results produced by Algorithm 7, it is essential for understanding the notion of datasets and sorting functions.

Initial vertex ranks in Algorithm 7 are computed by rounding the return value of r_{inf} to the next prime integer. Algorithm 8 shows how r_{inf} is actually implemented. The input to Algorithm 8 is a vertex of either a FCG or CG. The *if* condition, at lines 3–13, distinguishes between the two aforementioned cases; if v is a function (i.e. a vertex in a FCG),

Algorithm 8 Ordering function based on Shannon information

```

1: procedure  $r_{inf}(v)$ 
2:    $r \leftarrow 0$ 
3:   if  $v$  is function then
4:     for all  $basic\_block \in v$  do
5:       for all  $instruction \in basic\_block$  do
6:          $r \leftarrow r - \log p_{instruction}$ 
7:       end for
8:     end for
9:   else
10:    for all  $F \in v$  do
11:       $r \leftarrow r + r_{inf}(F)$ 
12:    end for
13:   end if
14:   return  $r$ 
15: end procedure

```

the return value is computed by summing $\log p_{instruction}$ for each instruction in the function's body (lines 4–8). Quantity $\log p_{instruction}$ is, in fact, the Shannon information content of an instruction computed on the random variable $X = P_1 \cup P_2$ i.e. the logarithm of the probability of occurrence of each instruction in the union of instructions of the programs under comparison which, practically, highlights functions that are composed of “rare” instructions. Returning back to the *if* clause, if v is a community of functions (or a community of communities), r_{inf} is called recursively for each element (lines 10–13). Finally, the computed value r is returned at line 14.

Algorithm 9 Ordering function based on vertex ranks

```

1: procedure  $r_{rank}(v)$ 
2:   return  $v.rank$ 
3: end procedure

```

Finally, r_{rank} , the sorting function used in Algorithm 6, can be seen in Algorithm 9. Since vertex ranks, assigned by `compute_vertex_ranks`, are stored as vertex attributes, r_{rank} returns that attribute value. This, effectively, means that, functions that process datasets, whose sorting function is set to r_{rank} , traverse graph vertices in increasing vertex ranks. This is especially true for functions `detect_initial_comms` and `improve_comms` used by Algorithm 6.

5.3.2 Algorithm 10: `compute_edge_ranks`

Once vertex ranks have been determined, `compute_edge_ranks` of Algorithm 10 is used to compute edge ranks. First, the PageRank [36] algorithm is used, at line 3, to compute initial rank values. We assume these values are stored in an attribute named *pagerank* of each edge of G . Then, at line 6, the rank of each edge (u, v) is set to a linear combination of

Algorithm 10 Compute initial edge ranks

```

1: procedure compute_edge_ranks( $D$ )
2:    $G \leftarrow G(D)$ 
3:    $\text{pagerank}(G)$ 
4:   for all  $e \in E(G)$  do
5:      $u, v \leftarrow e$ 
6:      $(u, v).\text{rank} \leftarrow \lceil u.\text{rank} * u.\text{pagerank} + v.\text{rank} * v.\text{pagerank} \rceil_p$ 
7:   end for
8: end procedure

```

the ranks of u and v , with the coefficients being the pagerank values of u and v respectively.

5.4 Algorithm 11: detect_initial_comms

Algorithm 11 shows the pseudocode for `detect_initial_comms`, our modified NSA community detection procedure. Recall that NSA is used to detect initial communities in each iteration of Algorithm 6.

Algorithm 11 NSA-based community detection

```

1: procedure detect_initial_comms( $D$ )
2:    $G \leftarrow G(D)$ 
3:    $\text{cnt} \leftarrow 0$ 
4:   for all  $v \in V(G)$  do
5:     if  $v.\text{comm}$  is set then
6:       continue
7:     end if
8:     if  $\text{succ}(v) = \emptyset$  and  $\text{pred}(v) = \emptyset$  then
9:        $v.\text{comm} \leftarrow 0$ 
10:      continue
11:    end if
12:     $u \leftarrow \text{max\_similarity\_neighbor}(v)$ 
13:    if  $u.\text{comm}$  is set then
14:       $v.\text{comm} \leftarrow u.\text{comm}$ 
15:    else
16:       $\text{cnt} \leftarrow \text{cnt} + 1$ 
17:       $v.\text{comm} \leftarrow \text{cnt}$ 
18:       $u.\text{comm} \leftarrow \text{cnt}$ 
19:    end if
20:  end for
21: end procedure

```

For each vertex (function or community of functions) in the dataset's graph (FCG or CG respectively) (line 4), it is first checked whether the vertex in question already belongs to a community (line 5). In this case, processing continues with the next vertex in the graph $G(D)$. It is crucial to mention that vertices of $G(D)$ are iterated based on the dataset's ordering function, which, in this case, happens to be r_{rank} (see how `detect_initial_comms` is invoked in Algorithm 6). Orphan vertices (i.e. vertices with no successors and predecessors) are always added in the same community, namely community 0 (lines 8–11), as this greatly simplifies the overall recovered community structure. The process

continues by considering the next vertex according to r_{rank} . For each v , the vertex most similar to v (as reported by `max_similarity_neighbor`, Algorithm 12, presented below) is returned in u (line 12) and the two are added in the same community (lines 13–14). However if $u.\text{comm}$ is not set (i.e. u does not belong to a community), a new community is created and v and u are both assigned to that community (lines 16–18).

Algorithm 12 Pick most similar neighbor of a vertex

```

1: procedure max_similarity_neighbor( $G, v$ )
2:    $S \leftarrow \text{argmax}_{u \in \text{neigh}(v)} \left( \frac{\|\text{neigh}(v) \cap \text{neigh}(u)\|}{\|\text{neigh}(v) \cup \text{neigh}(u)\|} \right)$ 
3:   if  $\|S\| > 1$  then
4:      $S \leftarrow \text{argmin}_{u \in S} (\|\text{neigh}(u)\|)$ 
5:     if  $\|S\| > 1$  then
6:        $S \leftarrow \text{argmin}_{u \in S} (\text{address}(v) - \text{address}(u))$ 
7:     end if
8:   end if
9:   return  $u \in S$ 
10: end procedure

```

Assigning vertices to communities is performed based on vertex similarity. The procedure that actually computes the similarity score between two vertices is given in Algorithm 12. For a vertex v , `max_similarity_neighbor` works as follows:

1. The ratio of common neighbors, between v and each of its neighbors u , over the union of their neighbors is computed (line 2). The neighbors of v with the highest ratio are added in set S and, if $\|S\| = 1$, the single element u of S is returned (line 9).
2. If $\|S\| > 1$, S is reduced to the set of those vertices with the smallest degree. If there are no ties (i.e. $\|S\| = 1$), u from S is returned (line 9).
3. If still $\|S\| > 1$, the vertex $u \in S$ whose address is nearest to v is picked. There can be no ties in this case. In case D represents a community dataset, `address()` can be made to return an arbitrary community characteristic (e.g. the lowest address member of the community).

5.5 Algorithm improve_comms

Algorithm 6 invokes `improve_comms` to fine-tune the partitioning returned by `detect_initial_comms`, using LMO heuristics. As it has already been mentioned, the former uses the directed Louvain algorithm. We will presently explain in detail the operation of this algorithm, but let us remark that it is the algorithm presented in [16], with a single modification: graph vertices are traversed in a deterministic way, based on the ordinals of r_{rank} , as opposed to random traversal proposed in the literature (refer to `improve_comms`, invoked in Algorithm 6).

For brevity purposes, we present our directed Louvain variant in verbal rather than pseudocode form. The algorithm can be summarized as follows:

1. Assign each vertex v of graph G in its own singleton community.
2. For each node v in G :
 - (a) Compute the potential increase in modularity by moving a neighbor u of v in the same community as v .
 - (b) Only move the vertex u that produces the maximum increase in modularity, if any.
3. Build induced graph of updated communities.
4. Repeat until modularity does not increase above a certain threshold.

The order, by which vertices in G are traversed in step 2, is arbitrary and might, as well, be random. This order, however, greatly affects the community structure recovered, and, so, for more deterministic results, a way of ordering nodes should be agreed upon before using Louvain to group functions in programs P_1 and P_2 . Sorting functions r_{rank} aim to partially tackle this problem, by assigning ranks to graph vertices and edges. Ranks are ordinals, which are used for sorting graph vertices and edges. As already mentioned, our Louvain variant traverses the graph vertices in that order, at step 2, as opposed to any other arbitrary order. This way, the more P_1 and P_2 “look alike”, the more their community structure, as recovered by our Louvain variant, does too.

6 LSH

6.1 Introductory remarks

Algorithm 5 delegates actual dataset matching to the following algorithms; `singleton_matcher`, `structural_matcher` and `monotonic_matcher`. The aforementioned procedures implement *matching strategies* and their semantics are discussed in [28]. In this work, we extend these semantics, by appending an additional parameter in their argument lists; a boolean value, named *is_inexact*, which signifies whether exact or inexact matching is to be performed.

In earlier versions of our work, to detect exact matches, a greedy $\mathcal{O}(nm)$ algorithm was used, where n was the size of the primary dataset and m that of the secondary. The algorithm iterated through both datasets and matched singleton elements (elements appearing only once in their respective datasets) with numerically identical feature vectors. This algorithm was, in turn, used to implement the three aforementioned matching strategies (singleton matching, structural

matching and monotonic matching). Greedy exact matching, which has been the standard practice in binary diffing tools, gives good results but performance drops dramatically as n and m increase. When inexact matching is also considered, $\mathcal{O}(nm)$ becomes prohibitively expensive, as distance computations between feature vectors cannot be assumed to take $\mathcal{O}(1)$ time (as assumed in the exact matching case).

Locality Sensitive Hashing (LSH) is a data clustering technique, which can remedy the greedy algorithm problems in the following ways:

1. *Dimensionality reduction*—LSH schemes can be used to approximate the Jaccard similarity of compared elements directly, by only considering the elements’ hash values. These hash values can be seen as the new set of the elements’ reduced dimensions.
2. *Bucketing*—LSH can aid in clustering data in smaller buckets, which can, then, be processed by divide and conquer algorithms to amortize the cost of $\mathcal{O}(nm)$.

REveal exploits the advantages offered by hashing as explained below:

1. In the inexact matching case, REveal implements a form of minhashing, based on *universal hash functions* [5] instead of permutations. Minhashing is used to split input datasets into smaller datasets consisting of elements that “look alike” according to their numerical features. Matching of elements between datasets is performed using the $\mathcal{O}(mn)$ greedy algorithm on the reduced datasets.
2. In the exact matching case, LSH does not offer any benefits. Instead plain hashing is used to split input datasets in smaller datasets consisting of elements with equal features. The new, smaller datasets are then processed by the $\mathcal{O}(mn)$ greedy algorithm.

6.2 Algorithm 13: `match_datasets`

The entry point of dataset matching logic is given in Algorithm 13.

Algorithm 13 Dataset matching entry point

```

1: procedure match_datasets( $D_1$ ,  $D_2$ ,  $M$ , is_inexact)
2:   if is_inexact = false then
3:      $n \leftarrow \text{match\_datasets\_exact}(D_1, D_2, M)$ 
4:   else
5:      $n \leftarrow \text{match\_datasets\_inexact}(D_1, D_2, M)$ 
6:   end if
7:   return  $n$ 
8: end procedure

```

At line 2, the value of *is_inexact* is checked. This boolean determines whether datasets will be matched exactly (line 3) or approximately (line 5).

6.2.1 Algorithm 14: *match_datasets_inexact*

Given two datasets D_1 and D_2 , Algorithm 14 begins by iterating through the vertices of graph $G(D_1)$, the primary dataset's graph, and hashing the feature vector of each vertex using all hash functions h_i in H (lines 3–8). All h_i are simple universal hash functions of the following form:

$$h_i(x) = (a_i x + b_i) \bmod p$$

Where p is a prime integer. Note that $G(D_1)$ can be either be a FCG or a CG, depending on the diffing phase, while H is assumed to be a set of randomly chosen universal hash functions. Next, for each k -shingle (k can be 1) w of the set of hashes S , a mapping $w \rightarrow v$ is created in the hash table T_1 (line 6) and, consequently, vertex v might be stored in several buckets of T_1 . The aforementioned process is then repeated for the graph of the secondary dataset $G(D_2)$ (lines 10–15). This time, hashed elements are added in hash table T_2 .

Algorithm 14 LSH-based inexact matching

```

1: procedure match_datasets_inexact( $D_1, D_2, M$ )
2:    $T_1 \leftarrow \emptyset$ 
3:   for all  $v \in G(D_1)$  do
4:      $S \leftarrow \{h_i(f(v)) \mid h_i \in H\}$ 
5:     for all  $w \in \text{shingles}(S)$  do
6:       insert( $T_1, w \rightarrow v$ )
7:     end for
8:   end for
9:    $T_2 \leftarrow \emptyset$ 
10:  for all  $v \in G(D_2)$  do
11:     $S \leftarrow \{h_i(f(v)) \mid h_i \in H\}$ 
12:    for all  $w \in \text{shingles}(S)$  do
13:      insert( $T_2, w \rightarrow v$ )
14:    end for
15:  end for
16:   $n \leftarrow 0$ 
17:  for all  $w \in \text{keys}(T_1) \cap \text{keys}(T_2)$  do
18:     $D'_1 \leftarrow \langle G(D_1), \text{values}(T_1, w), r(D_1), f(D_1) \rangle$ 
19:     $D'_2 \leftarrow \langle G(D_2), \text{values}(T_2, w), r(D_2), f(D_2) \rangle$ 
20:     $n \leftarrow n + \text{greedy}(D'_1, D'_2, M)$ 
21:  end for
22:  return  $n$ 
23: end procedure

```

In the following loop (lines 17–21), only the set of keys that both T_1 and T_2 have in common are considered. For each such key w , a temporary dataset D'_1 , consisting of D_1 's vertices in $T_1[w]$ (line 18), and a temporary dataset D'_2 , consisting of D_2 's vertices in $T_2[w]$ (line 19), are constructed and a greedy matching algorithm (function *greedy*) is invoked to solve the resulting reduced problem. We assume that given

two datasets, *greedy* produces a matching of vertices with a low (but not necessarily the lowest) cost, adds matched pairs in M , removes matched vertices from their respective datasets and returns the number of matches n .

6.2.2 Algorithm 15: *match_datasets_exact*

For exact matching purposes, *match_datasets_exact* uses an arbitrary hash function h . This can be either a cryptographic hash function, or even a CRC code computed over $f(v)$, the feature vector of a vertex v . This is done for all v in datasets D_1 (lines 3–6) and D_2 (lines 8–11) with hashed elements ending up in hash tables T_1 and T_2 respectively. The loop at lines 13–17 is similar to that of Algorithm 14. We again assume that *greedy* populates M with the newly matched pairs, removes vertices from their datasets and returns the number n of matched pairs.

Algorithm 15 Hashing-based exact matching

```

1: procedure match_datasets_exact( $D_1, D_2, M$ )
2:    $T_1 \leftarrow \emptyset$ 
3:   for all  $v \in G(D_1)$  do
4:      $w \leftarrow h(f(v))$ 
5:     insert( $T_1, w \rightarrow v$ )
6:   end for
7:    $T_2 \leftarrow \emptyset$ 
8:   for all  $v \in G(D_2)$  do
9:      $w \leftarrow h(f(v))$ 
10:    insert( $T_2, w \rightarrow v$ )
11:  end for
12:   $n \leftarrow 0$ 
13:  for all  $w \in \text{keys}(T_1) \cap \text{keys}(T_2)$  do
14:     $D'_1 \leftarrow \langle G(D_1), \text{values}(T_1, w), r(D_1), f(D_1) \rangle$ 
15:     $D'_2 \leftarrow \langle G(D_2), \text{values}(T_2, w), r(D_2), f(D_2) \rangle$ 
16:     $n \leftarrow n + \text{greedy}(D'_1, D'_2, M)$ 
17:  end for
18:  return  $n$ 
19: end procedure

```

7 Experimental results

7.1 Baseline

In this section we compare REveal, our binary diffing framework, against two popular tools of the same domain, namely, Diaphora and YaDiff. All three run on IDA Pro; the first two as IDA Python plug-ins, while the latter as a native plug-in. Furthermore, for evaluating the efficiency of CD and LSH, we also compare the current version of REveal against its older version that lacks inexact matching capabilities and the algorithm implementations presented in the previous sections.

REveal is the name of our binary diffing software, first introduced in [27]. It implements all algorithms presented

in [27,28], as well as the ones of the present paper. As our research progresses, we extend it with new algorithms and features, constantly improving its matching efficiency and speed. Even though REveal is not yet in the public domain, we plan to release it once its code base matures and becomes stable. For the time being, access, for non-commercial use and for preview purposes, can be granted, on request, to interested researchers and reverse engineers.

Diaphora is a free and open-source binary diffing tool, which is probably the most widespread of its kind in the reverse engineering community. It was first released during SyScan 2015 and is still actively maintained. Diaphora works by exporting function information, from IDA Pro databases under comparison, into portable SQLite files. This per-function information includes the number of basic blocks, the number of connected components, a signature synthesized by the function's instructions and many more. Diaphora is also capable of utilizing the power offered by the IDA decompiler. If the latter is detected, Diaphora uses it to create a filtered version of each function's decompiled C code and stores the result as yet another function feature. Diaphora compares the exported databases using a set of predetermined SQL queries, which compare various combinations of function features. These queries are able to detect both exact, as well as inexact matches. In the latter case, a similarity score is also returned. It is worth noting that Diaphora also comes with a graphical user interface that allows reverse engineers to explore the diffing results interactively.

YaDiff is a binary diffing tool, part of YaCo, a collaborative reverse engineering plug-in for IDA Pro. YaDiff was developed for the purpose of porting information between two IDA Pro databases. That is, given two IDBs, it applies a binary diffing algorithm for the purpose of merging information, from the former to the latter, and writes the result in a new database, so as not to trash the original inputs. Merged information includes, symbol names, comments, bookmarks and all sorts of annotations usually found in an IDA Pro database. That said, YaDiff is not a generic, interactive binary diffing tool (like Diaphora which allows users to examine the list of symbols that were matched between the two databases), it merely performs the information merging process in a fully automated manner, giving no insights on the actual matching results. However, since YaDiff is open source and function matching is an integral part of its binary diffing engine, we were able to modify it and make it expose the required information. With the said modifications in place, YaDiff's matching power is now measurable.

7.2 Evaluation

All experiments were performed on IDA Pro 7, running IDA Python based on Python 2.7. For this purpose, the version of Diaphora used is the latest from the *diaphora-1.2* branch,

with a series of patches manually backported from the *master* branch, to fix minor bugs in the former. Unfortunately, running the latest version of Diaphora, requires a new IDA Pro license, as both have recently moved to using Python 3 instead. When it comes to YaDiff, the latest commits in its *master* branch appear to be about 3 years old. For the purpose of running our experiments, the latest version of YaDiff was modified, as already mentioned in the previous paragraph, and was manually compiled using the IDA Pro SDK. Last but not least, it should be noted that all experiment instances were carried out on a Mac laptop equipped with a Core i7 2.4Ghz and 8Gb of RAM. We plan to perform comparisons against the latest Diaphora version (as well as against additional binary diffing tools) in the near future and report them in another publication.

In the following, all binary diffing tools were used to find both exact and inexact function matches in the compared subjects. Debugging information, present in the experiment corpus, is used as a ground truth for verifying the correctness of the results and classifying them as either valid matches or mismatches. For the sake of fairness, function names in IDA Pro databases were scrambled (using simple IDA Python scripts) before executing the diffing engines. Note that the time it takes for the compared tools to export information is not measured; only the diffing process is timed and displayed in the following tables.

To demonstrate the efficiency of our new, divide-and-conquer approach in binary diffing, as well as the improvements in REveal's diffing engine for that matter, we reused the same dataset as in [28]. The dataset consists of binary executables for *ncmc*, *nmap*, *ffmpeg-android* and the Linux kernel. The aforementioned programs are of increasing complexity, with *ncmc* executables having the least number of functions (≈ 1000) and *vmlinux* binaries the most (≈ 40000). The dataset contains two versions of each of the aforementioned programs, for example, Linux kernel 4.4.1 and 4.4.40, which are diffed with one another. Furthermore, for each pair of versions of each program, executables for all four computer architectures, namely *i386*, *amd64*, *arm* and *aarch64* are compared. The above combinations result in a total of 16 experiment instances summarized in Fig. 2. For each experiment, we measure the time it takes for the diffing process to complete (ignoring latencies introduced during information exporting), the number of successful function matches, as well as the number of mismatches, that is, functions that were matched, but according to ground truth, they are not the same.

Additionally, we evaluate the abovementioned diffing tools in a subset of the DeepBinDiff [13] dataset, which is publicly available at the project's GitHub repository. This dataset contains several precompiled binary executables, generated by compiling various versions of *coreutils*, *diffutils* and *findutils*, using optimization levels ranging from

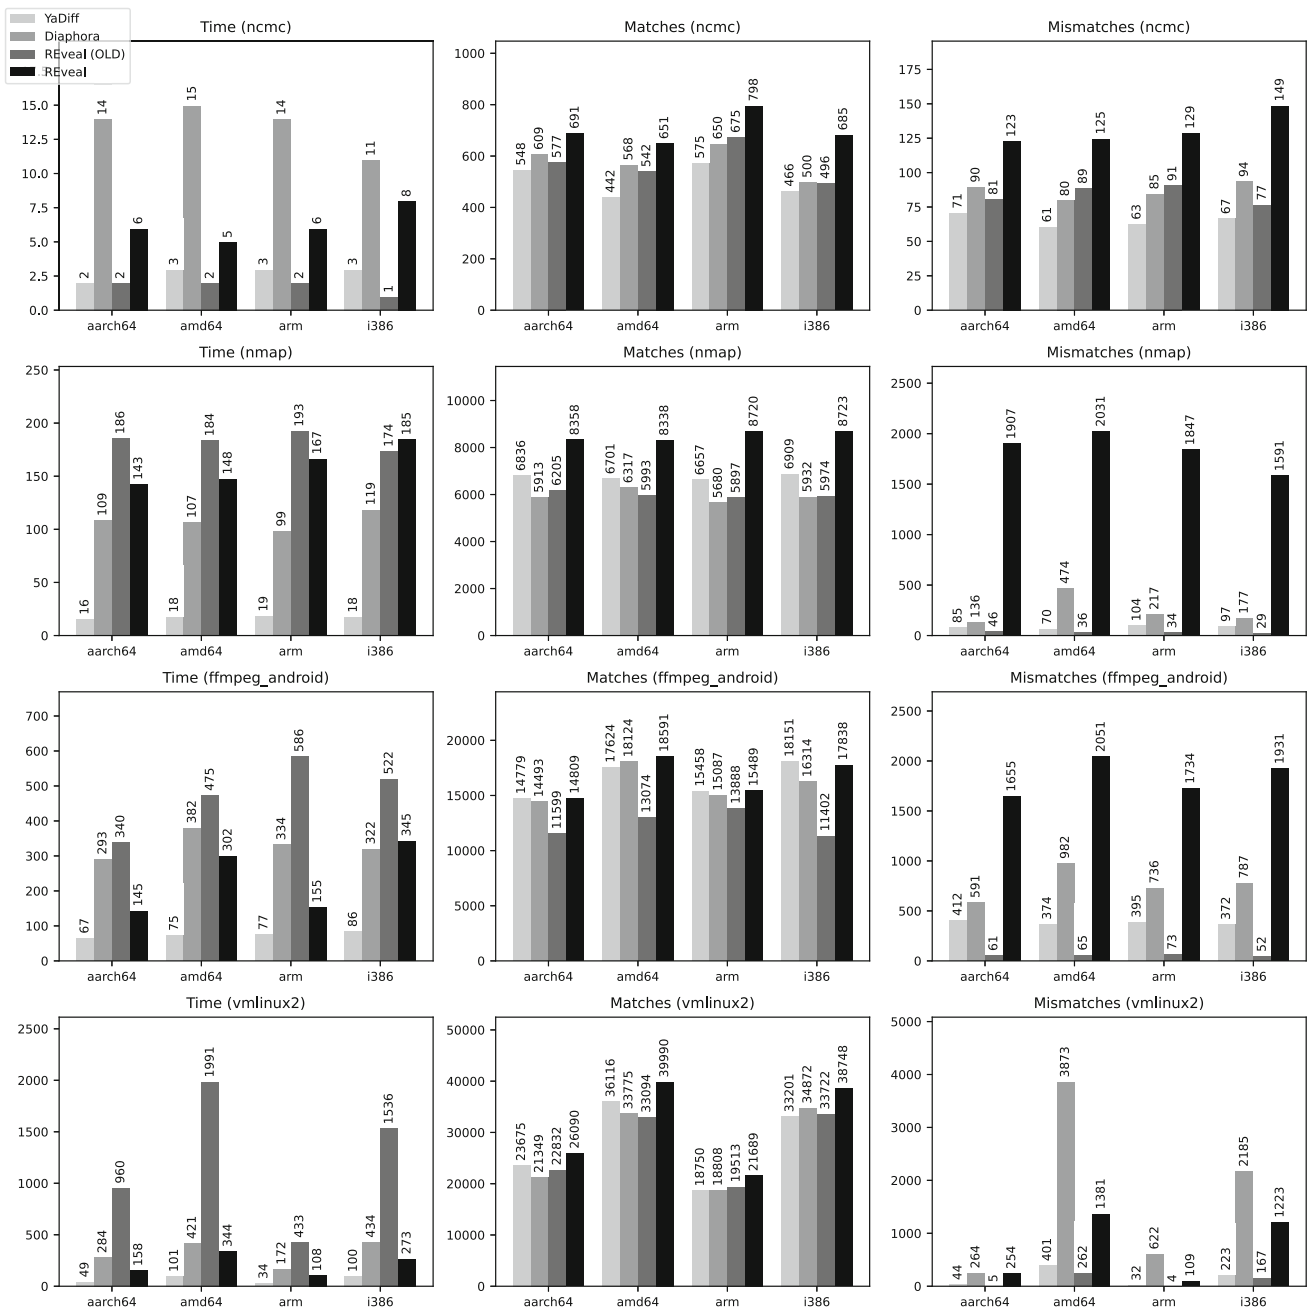


Fig. 2 Experimental results for our custom dataset. Rows correspond to programs (*ncmc*, *nmap*, *ffmpeg-android*, *vmlinux*) and columns correspond to measurements (execution time, number of matches, number of mismatches). For the total number of functions in each case refer to Table 1

O0 to O3. More specifically, during our evaluation, we considered the following; *coreutils* version 8.1 against 8.30 (102 binaries), *diffutils* version 3.4 against 3.6 (4 binaries) and *findutils* version 4.41 against 4.6 (3 binaries). Executables, compiled with optimization levels from O0 to O3, were used, resulting in a total of 12 experiment instances summarized in Fig. 3. For each experiment, we take the exact same measurements as previously (execution time, number of matches and

mismatches). Due to lack of historical data, the performance of REveal's predecessor is not shown. Note that Fig. 3 shows only the accumulated execution time, number of matches and mismatches, not the individual diffing results of each diffed binary executable pair.

7.3 Results

7.3.1 Custom dataset

Experimental results are shown in Fig. 2 (see also Table 1 for the total number of functions contained in each program and architecture). We observe that REveal has the highest number of correctly matched functions in all experiment instances apart from one (*ffmpeg_android* for *i386*). Starting with *ncmc* (Fig. 2, 1st row), REveal successfully matches 691, 651, 798 and 685 functions, when comparing *aarch64*, *amd64*, *arm* and *i386* binary executables respectively. Diaphora comes second with 609, 568, 650 and 500 matches. The matching power of REveal, compared to this of its previous version, presented in [28], has evidently risen, at the expense of slightly increased running time. YaDiff comes third with respect to number of matches, but has shorter running times than its two competitors. REveal produces the most false positives, with an average of about 132 mismatches in each instance of this first round of experiments.

The second series of experiments (Fig. 2, 2nd row) concerns *nmap*. The subject binaries, in this batch of experiments, are larger, with the number of functions being increased from a few hundreds, compared to *ncmc*, to a few thousands. Once again we see a higher number of successful matches detected by REveal, which beats YaDiff by 1522, 1637, 2063 and 1814 matches. Interestingly, in the first three

instances, the updated version of REveal finds a few thousand extra matches, compared to its previous version, about 10–40 s faster. As the size of the compared programs increases, YaDiff’s execution time advantage becomes more marked, with its running times being an order of magnitude shorter than those of its competitors. It’s also obvious that, in this series of experiment, YaDiff outperforms Diaphora in terms of correctly matched functions. Last but not least, when it comes to mismatches, REveal is again the loser, with an average of about 1844 mismatches per run.

The next four experiment instances (Fig. 2, 3rd row) concern *ffmpeg* binaries for Android. The number of functions have now doubled, compared to *nmap*. We begin by noting that, for *i386*, YaDiff outperforms both Diaphora and REveal, with a score of 18151 matches and a running time of 86 s. The new version of REveal is both faster and more efficient, in terms of matches, than its predecessor and is the winner in the remaining runs. When it comes to mismatches, REveal is again the loser, but the percentages of erroneous matches have fallen to more tolerable values, compared to the total number of functions. YaDiff performs exceptionally well, even outperforming Diaphora in some runs, with impressive speed efficiency, making it ideal for quick binary diffing tasks.

The *vmlinux* series of experiments (Fig. 2, 4th row) highlight REveal’s power to match binary executables, when these are very similar. First and foremost, we can see that REveal is once again the winner when looking at the numbers of correctly matched functions. Additionally, even though the number of mismatches, found by REveal, is evidently higher than this of YaDiff, we can see that, generally, it’s Diaphora that makes the most mistakes in this batch. In terms of execution time, YaDiff is followed by REveal, which in turn wins Diaphora, while the obsolete version of REveal comes last. Evidently REveal has greatly benefited by CD and LSH, when it comes to execution speed compared to its predecessor. The second and fourth experiment instances highlight this, with the new version executing about 6 to 7 times faster.

Table 1 Number of functions in our custom dataset

Program	Arch.	V. 1	V. 2
ncmc		0.1.7	0.1.8
	aarch64	820	822
	amd64	780	782
	arm	931	933
	i386	839	837
nmap		7.12	7.31
	aarch64	10773	13344
	amd64	10806	13374
	arm	11526	13469
	i386	10923	13504
ffmpeg-android		20180408	20180731
	aarch64	16489	20870
	amd64	20649	25028
	arm	17226	21610
	i386	19800	24174
vmlinux		4.4.1	4.4.40
	aarch64	26383	26430
	amd64	41598	41527
	arm	21813	21852
	i386	40111	40184

7.3.2 DeepBinDiff dataset

As previously mentioned, the DeepBinDiff dataset consists of several ELF binary executables which are, generally, small in size (as opposed to our custom dataset that mainly consists of a few large executables). One conclusion, that can be quickly drawn by looking at Fig. 3 (and Table 2 that shows the total number of functions for each program and optimization level), is that, REveal is about 2 to 5 times slower when compared to its competitors. CD and LSH both require “setup” times, whose cost is amortized for larger executables, but becomes more evident when the compared subjects are smaller. As it can be seen, in all experiment instances (Fig. 3, 1st column) REveal takes 17 to 364 s to complete,

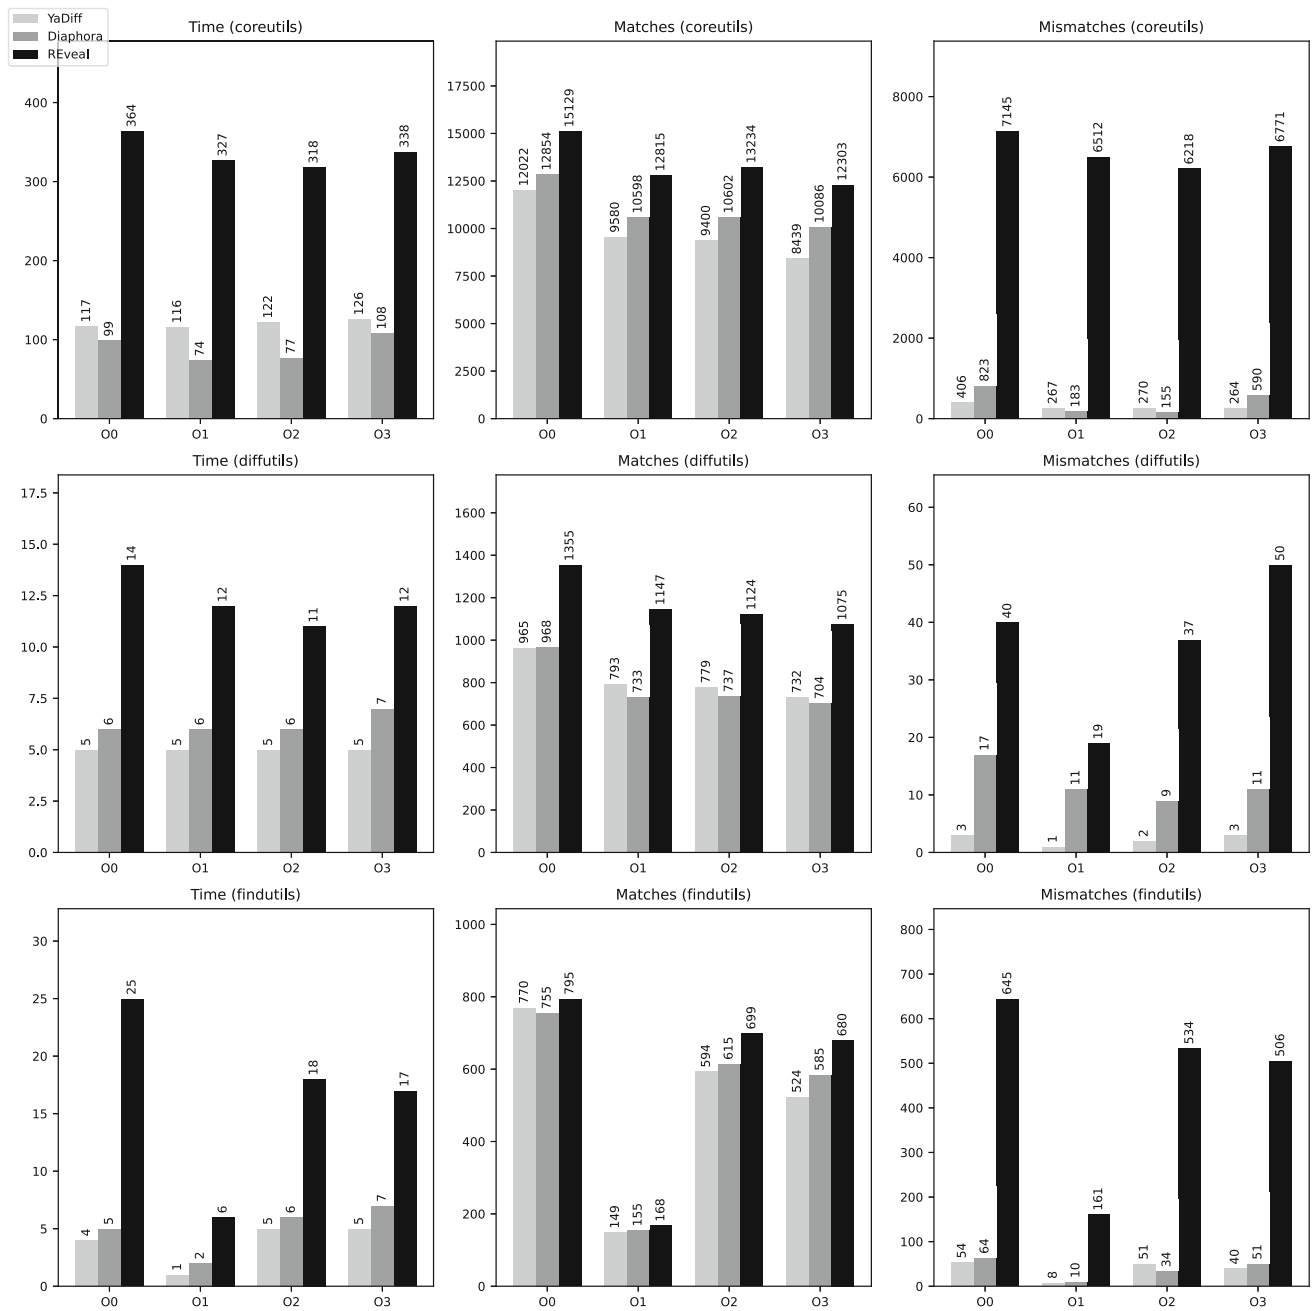


Fig. 3 Experimental results for the DeepBinDiff dataset. Rows correspond to programs (*coreutils*, *diffutils*, *findutils*) and columns correspond to measurements (execution time, number of matches, number of mismatches). For the total number of functions in each case refer to Table 2

while both Diaphora and YaDiff are much faster, with running times ranging from 2 to 108 s for the former and from 1 to 126 s for the latter.

Focusing on the number of correctly matched functions (Fig. 3, 2nd column), we can see how REveal wins the race, with more noticeable leverage in the cases of *coreutils* and *diffutils* (rows 1 and 2 respectively). Of course, this additional precision comes with the time penalty mentioned in the previous paragraph. Diaphora and YaDiff also perform

well, but come second and third respectively, with the former winning the latter in 8 out of the 12 experiment instances (more specifically YaDiff wins in *diffutils* O1, O2 and O3, as well as in *findutils* O0).

Finally, when it comes to mismatches (Fig. 3, 3rd column), we again experience high numbers of false positives generated by REveal, especially in the *coreutils* and *findutils* cases, with the O0 and O2 *findutils* experiments being the most notable ones. In these two cases, REveal generates a

Table 2 Number of functions in DeepBinDiff dataset (for reasons unknown to us, *findutils* O1 consists of a single binary, hence the inconsistency in the numbers)

Program	Optimization level	V. 1	V. 2
coreutils		8.1	8.30
	O0	23503	28485
	O1	20096	23534
	O2	20178	23453
	O3	19685	22853
diffutils		3.4	3.6
	O0	1408	1408
	O1	1171	1173
	O2	1168	1170
	O3	1129	1133
findutils		4.41	4.6
	O0	1663	1685
	O1	404	334
	O2	1356	1412
	O3	1250	1342

number of false positives almost equal to the number of true positives. On the contrary, Diaphora and YaDiff are more resilient and generate acceptable false positive percentages. This time, YaDiff wins by generating the least mismatches in 9 out of 12 runs, with Diaphora winning the O1 and O2 *coreutils* run, as well as the O2 *findutils* one.

7.4 On the number of false positives

From Figs. 2 and 3, and the analysis of the previous section, it is clear that REveal has a high false positive rate. During its inexact matching rounds, REveal will keep matching as many functions as it can, even if the similarity score is very low (i.e., distance is high), as long as unique score matches can be found and, consequently, this may lead to increased numbers of false positives. Trying to avoid this by hardcoding distance thresholds, below which matches are considered legitimate, and dropping all matches above that threshold, is generally not an acceptable solution.

However, the high false positive rate is not an acute problem. Among others, binary diffing tools aid in manual reverse engineering and REveal is not an exception. Since all results are provided to the user, matches can be filtered using application-specific and domain-specific knowledge at a later time. Indeed, as a reverse engineer makes use of the results produced by a diffing tool, he or she can decide whether a match has been identified correctly or not. Based on manual effort, user-identified matches can be specified and guide the overall matching process. Furthermore, research on more reliable community detection and community matching algo-

rithms that will allow to, at least partially, tackle this problem, is currently underway.

8 Conclusion

We have empirically proven that binary diffing can benefit from divide-and-conquer approaches, using techniques like CD and LSH, both in terms of speed efficiency as well as matching precision. However, we believe this merits further research. Dataset partitioning during binary diffing is still an unpopular approach, so we hope to have inspired researchers and reverse engineers to look into it. Several community detection and, generally, graph partitioning algorithms can be found in the public literature; their systematic testing and formal evaluation in binary diffing applications will definitely advance the current state of the art.

Funding Open access funding provided by HEAL-Link Greece

Declarations

Conflict of interest All authors certify that they have no affiliations with or involvement in any organization or entity with any financial interest or non-financial interest in the subject matter or materials discussed in this manuscript.

Open Access This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by/4.0/>.

References

1. Babai, L.: Graph isomorphism in quasipolynomial time. In: Proceedings of the Forty-Eighth Annual ACM Symposium on Theory of Computing, pp. 684–697 (2016)
2. BinDiff, Z.: <https://www.zynamics.com/bindiff.html>
3. Blondel, V.D., et al.: Fast unfolding of communities in large networks. *J. Stat. Mech. Theory Exp.* **10**, 1008 (2008)
4. Broder, A.Z.: On the resemblance and containment of documents. In: Proceedings of Compression and Complexity of Sequences (CCCS'97), pp. 21–29 (1997)
5. Carter, J.L., Wegman, M.N.: Universal classes of hash functions. *J. Comput. Syst. Sci.* **18**, 143–154 (1979)
6. Cesare, S., Xiang, Y.: Classification of malware using structured control flow. In: Proceedings of the 8th Australasian Symposium on Parallel and Distributed Computing (AusPDC 2010), vol. 107, pp. 61–70 (2010)

7. Cesare, S., Xiang, Y., Zhou, W.: Control flow-based malware variant detection. *IEEE Trans. Depend. Secure Comput.* **11**, 307–317 (2013)
8. Charikar, M.S.: Similarity estimation techniques from rounding algorithms. In: *Proceedings of the 34th Annual ACM Symposium on Theory of Computing*, pp. 380–388 (2002)
9. Cheng, J., et al.: Neighbor similarity based agglomerative method for community detection in networks. *Complexity* **1**, 1–16 (2019)
10. Cordeiro, M., Sarmento, R.P., Gama, J.: Dynamic community detection in evolving networks using locality modularity optimization. *Soc. Netw. Anal. Min.* **6**, 1–20 (2016)
11. Diaphora, J.K.: A Free and Open Source Program Diffing Tool. <http://diaphora.re/>
12. Dongkwan, K., Eunsoo, K., Sang, K. C., Soeul, S., Yongdae, K.: Revisiting Binary Code Similarity Analysis using Interpretable Feature Engineering and Lessons Learned, *Computing Research Repository (CoRR)*, abs/2011.10749 (2020)
13. Duan, Y., Li, X., Wang, J., Yin, H.: DeepBinDiff: Learning Program-Wide Code Representations for Binary Diffing, *NDSS* (2020)
14. Dullien, T., Rolles, R.: Graph-based comparison of executable objects. In: *Proceedings of the Symposium sur la Securite des Technologies de l'Information et des Communications*, pp. 3 (2005)
15. Dullien, T., Carrera, E., Eppler, S.M., Porst, S.: Automated attacker correlation for malicious code. In: *NATO Information Systems Technology (IST) 091*, pp. 26.1–26.10 (2010)
16. Dugue, N., Perez, A.: Directed Louvain: Maximizing Modularity in Directed Networks. *Universite d'Orleans, Research Report* (2015)
17. Flake, H.: Structural comparison of executable objects. In: *Proceedings of the IEEE Conference on Detection of Intrusions and Malware and Vulnerability Assessment (DIMVA)*, pp. 161–173 (2004)
18. Fortunato, S., Barthelemy, M.: Resolution limit in community detection. *Proc. Natl. Acad. Sci.* **104**, 36–41 (2007)
19. Fortunato, S.: Community detection in graphs. *Phys. Rep.* **486**, 75–174 (2010)
20. Fortunato, S., Hric, D.: Community detection in networks: a user guide. *Phys. Rep.* **659**, 1–44 (2016)
21. Frago, L.D.S.: STABLE: Making Player Modeling Possible for Sandbox Games. *University of Saskatchewan, Ph.D. Diss* (2020)
22. Gionis, A., Indyk, P., Motwani, R.: Similarity search in high dimensions via hashing. In: *Proceeding VLDB '99 Proceedings of the 25th International Conference on Very Large Data Bases*, pp. 518–529 (1999)
23. Hu, X., Chiueh, T., Shin, K.G.: Large-Scale Malware Indexing Using Function-Call Graphs, *Computer and Communications Security*, pp. 611–620 (2009)
24. Jaja, J., Kosaraju, S.R.: Parallel algorithms for planar graph isomorphism and related problems. *IEEE Trans. Circuits Syst.* **35**, 304–311 (1988)
25. Indyk, P., Motwani, R.: Approximate nearest neighbors: towards removing the curse of dimensionality. In: *Proceedings of the Thirtieth Annual ACM Symposium on Theory of Computing*, pp. 604–613 (1998)
26. Jurczyk, M.: Using Binary Diffing to Discover Windows Kernel Memory Disclosure Bugs. <https://googleprojectzero.blogspot.gr/2017/10/using-binary-diffing-to-discover.html>
27. Karamitas, C., Kehagias, A.: Efficient features for function matching between binary executables. In: *IEEE 25th International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pp. 335–345 (2018)
28. Karamitas, C., Kehagias, A.: Function matching between binary executables: efficient algorithms and features. *J. Comput. Virol. Hack. Tech.* **15**, 307–323 (2019)
29. Karande, V., et al.: BCD: Decomposing binary code into components using graph-based clustering. In: *Proceedings of the 13th Asia Conference on Computer and Communications Security (AsiaCCS)*, pp. 393–398 (2018)
30. Kim, J., Kim, T.G., Im, E.G.: Structural information based malicious app similarity calculation and clustering. In: *Proceedings of the 2015 Conference on Research in Adaptive and Convergent Systems (RACS)*, pp. 314–318 (2015)
31. Kinable, J., Kostakis, O.: Malware classification based on call graph clustering. *J. Comput. Virol.* **7**, 233–245 (2011)
32. Lu, C., et al.: Graph Iso/auto-morphism: a divide-and-conquer approach. In: *Proceedings of the 2021 International Conference on Management of Data*, pp. 1195–1207 (2021)
33. Luo, L., Ming, J., Wu, D., Liu, P., Zhu, S.: Semantics-based obfuscation-resilient binary code similarity comparison with applications to software plagiarism detection. In: *Proceedings of the 22nd ACM SIGSOFT International Symposium on Foundations of Software Engineering*, pp. 389–400 (2014)
34. Ma, Y., Zhao, Y., Wang, J.: Modularity-based incremental label propagation algorithm for community detection. *Appl. Sci.* **10**, 4060 (2020)
35. Mengin, E.: Weisfeiler–Lehman Graph Kernel for Binary Function Analysis. <https://blog.quarkslab.com/weisfeiler-lehman-graph-kernel-for-binary-function-analysis.html>
36. Page, L., Brin, S., Motwani, R., Winograd, T.: The PageRank citation ranking: bringing order to the web. *Proc. ASIS* **98**, 161–172 (1999)
37. Qu, Y., et al.: Exploring community structure of software call graph and its applications in class cohesion measurement. *J. Syst. Softw.* **108**, 193–210 (2015)
38. Traag, V.A., Waltman, L., van Eck, N.J.: From Louvain to Leiden: guaranteeing well-connected communities. *Sci. Rep.* **9**, 1–12 (2019)
39. Ullah, S., Jin, W., Oh, H.: Efficient features for function matching in multi-architecture binary executables. *IEEE Access* **9**, 104950–104968 (2021)
40. Wang, Z., Pierce, K., McFarling, S.: BMAT—a binary matching tool. *J. Instruct. Level Parall.* **2**, 1–20 (2000)
41. Weisfeiler, B.Y., Lehman, A.A.: A reduction of a graph to a canonical form and an algebra arising during this reduction. In: *Nauchno Technicheskaya Informatsia*, vol. 9, pp. 12–16 (1968)
42. Xiu, H., et al.: Hierarchical Graph Matching Network for Graph Similarity Computation, *Computing Research Repository (CoRR)*, abs/2006.16551 (2020)
43. YaDiff, B.A., et al.: <https://github.com/DGA-MI-SSI/YaCo>

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.